

Arrakis: The Operation System is the control plane Reading Report

Simon Peter, University of Washionton .osdi

邓志会 2015210926

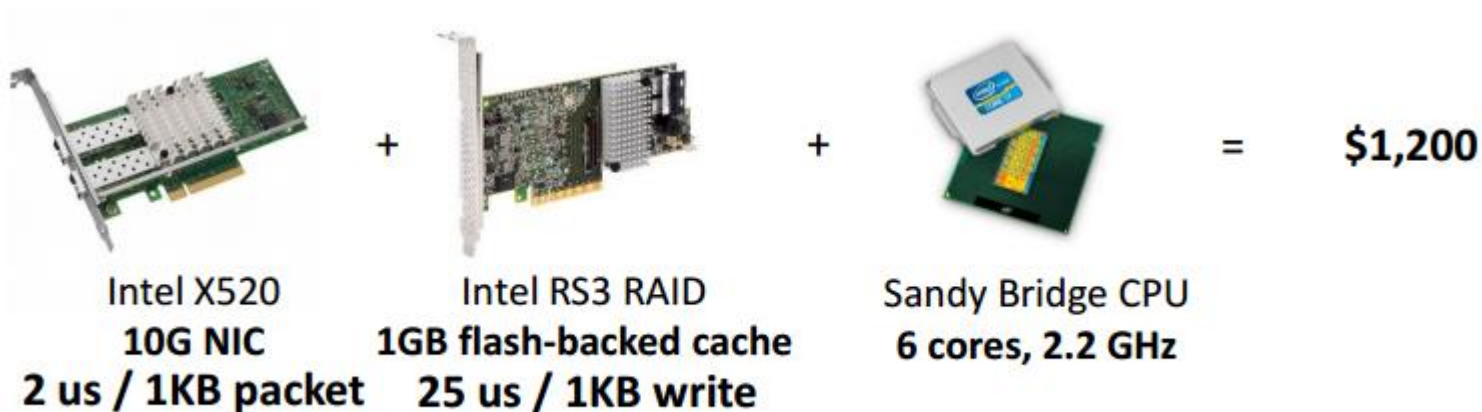
元东 2015210938

Building an os for data center

Server I/O performance matters

key-value stores, web & file servers, lock management

- Can we deliver performance close to hardware?
- Example system: Dell PowerEdge R520



Intel X520
10G NIC
2 us / 1KB packet

+

Intel RS3 RAID
1GB flash-backed cache
25 us / 1KB write

+

Sandy Bridge CPU
6 cores, 2.2 GHz

= \$1,200

Packet processing overhead

Table I. Sources of Packet Processing Overhead in Linux and Arrakis

All times are averages over 1,000 samples, given in μs (and standard deviation for totals). Arrakis/P uses the POSIX interface; Arrakis/N uses the native Arrakis Interface.

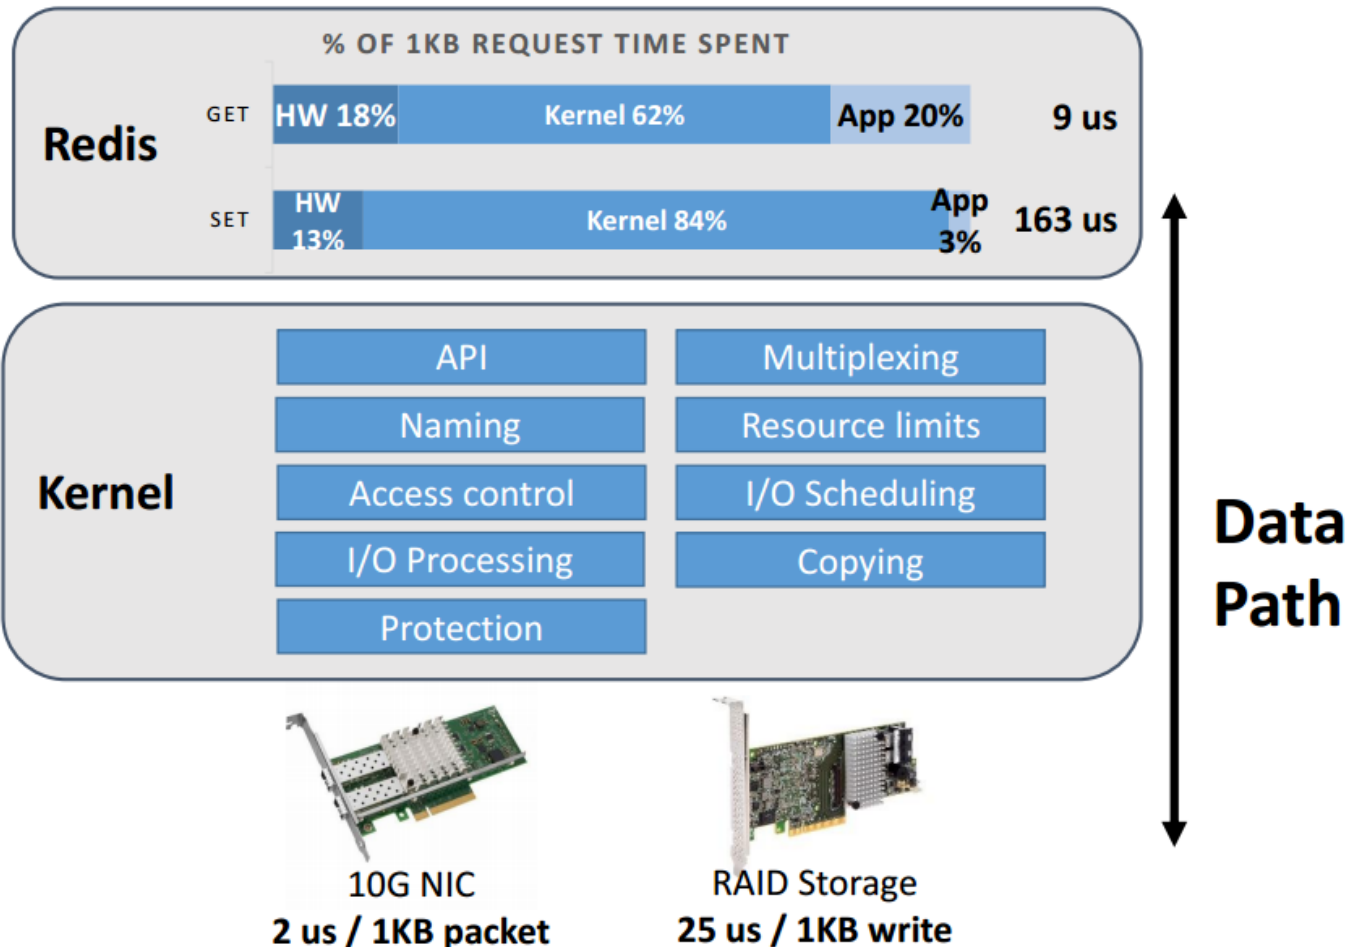
		Linux				Arrakis			
		Receiver running		CPU idle		Arrakis/P		Arrakis/N	
Network stack	in	1.26	(37.6%)	1.24	(20.0%)	0.32	(22.3%)	0.21	(55.3%)
	out	1.05	(31.3%)	1.42	(22.9%)	0.27	(18.7%)	0.17	(44.7%)
Scheduler		0.17	(5.0%)	2.40	(38.8%)	-		-	
Copy	in	0.24	(7.1%)	0.25	(4.0%)	0.27	(18.7%)	-	
	out	0.44	(13.2%)	0.55	(8.9%)	0.58	(40.3%)	-	
Kernel crossing	return	0.10	(2.9%)	0.20	(3.3%)	-		-	
	syscall	0.10	(2.9%)	0.13	(2.1%)	-		-	
Total		3.36		6.19		1.44		0.38	
Std. dev.		0.66		0.82		<0.01		<0.01	

Redis NoSqlstore overheads

	Read hit				Durable write			
	Linux		Arrakis/P		Linux		Arrakis/P	
epoll	2.42	(27.91%)	1.12	(27.52%)	2.64	(1.62%)	1.49	(4.73%)
recv	0.98	(11.30%)	0.29	(7.13%)	1.55	(0.95%)	0.66	(2.09%)
Parse input	0.85	(9.80%)	0.66	(16.22%)	2.34	(1.43%)	1.19	(3.78%)
Lookup/set key	0.10	(1.15%)	0.10	(2.46%)	1.03	(0.63%)	0.43	(1.36%)
Log marshaling	-		-		3.64	(2.23%)	2.43	(7.71%)
write	-		-		6.33	(3.88%)	0.10	(0.32%)
fsync	-		-		137.84	(84.49%)	24.26	(76.99%)
Prepare response	0.60	(6.92%)	0.64	(15.72%)	0.59	(0.36%)	0.10	(0.32%)
send	3.17	(36.56%)	0.71	(17.44%)	5.06	(3.10%)	0.33	(1.05%)
Other	0.55	(6.34%)	0.46	(11.30%)	2.12	(1.30%)	0.52	(1.65%)
Total	8.67	($\sigma=2.55$)	4.07	($\sigma=0.44$)	163.14	($\sigma=13.68$)	31.51	($\sigma=1.91$)
99th percentile	15.21		4.25		188.67		35.76	

Table 2: Overheads in the Redis NoSQL store for memory reads (hits) and durable writes (legend in Table 1).

Linux I/O Performance



Kernel mediation is too heavyweight

Arrakis Goals

Skip kernel & deliver I/O directly to applications

- Reduce OS overhead

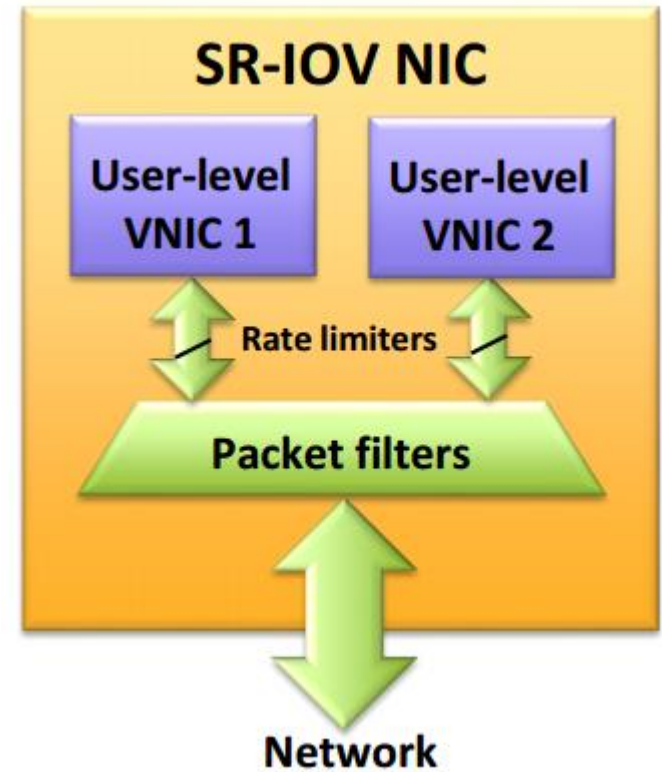
Keep classical server OS features

- Process protection
- Resource limits
- I/O protocol flexibility
- Global naming

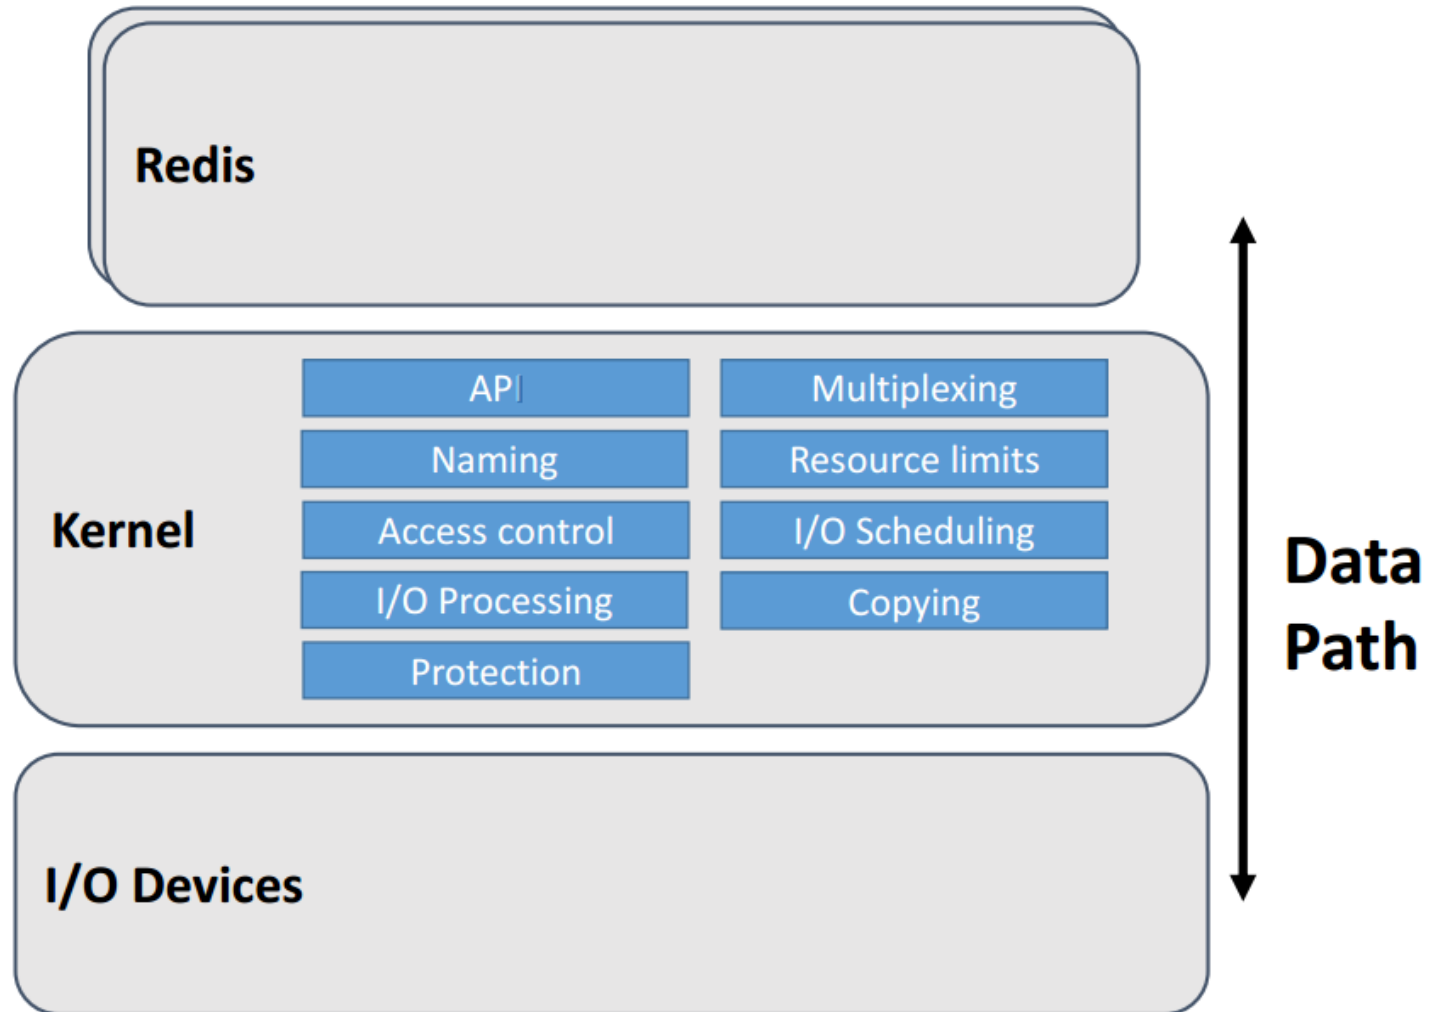
The hardware can help us...

Hardware I/O Virtualization

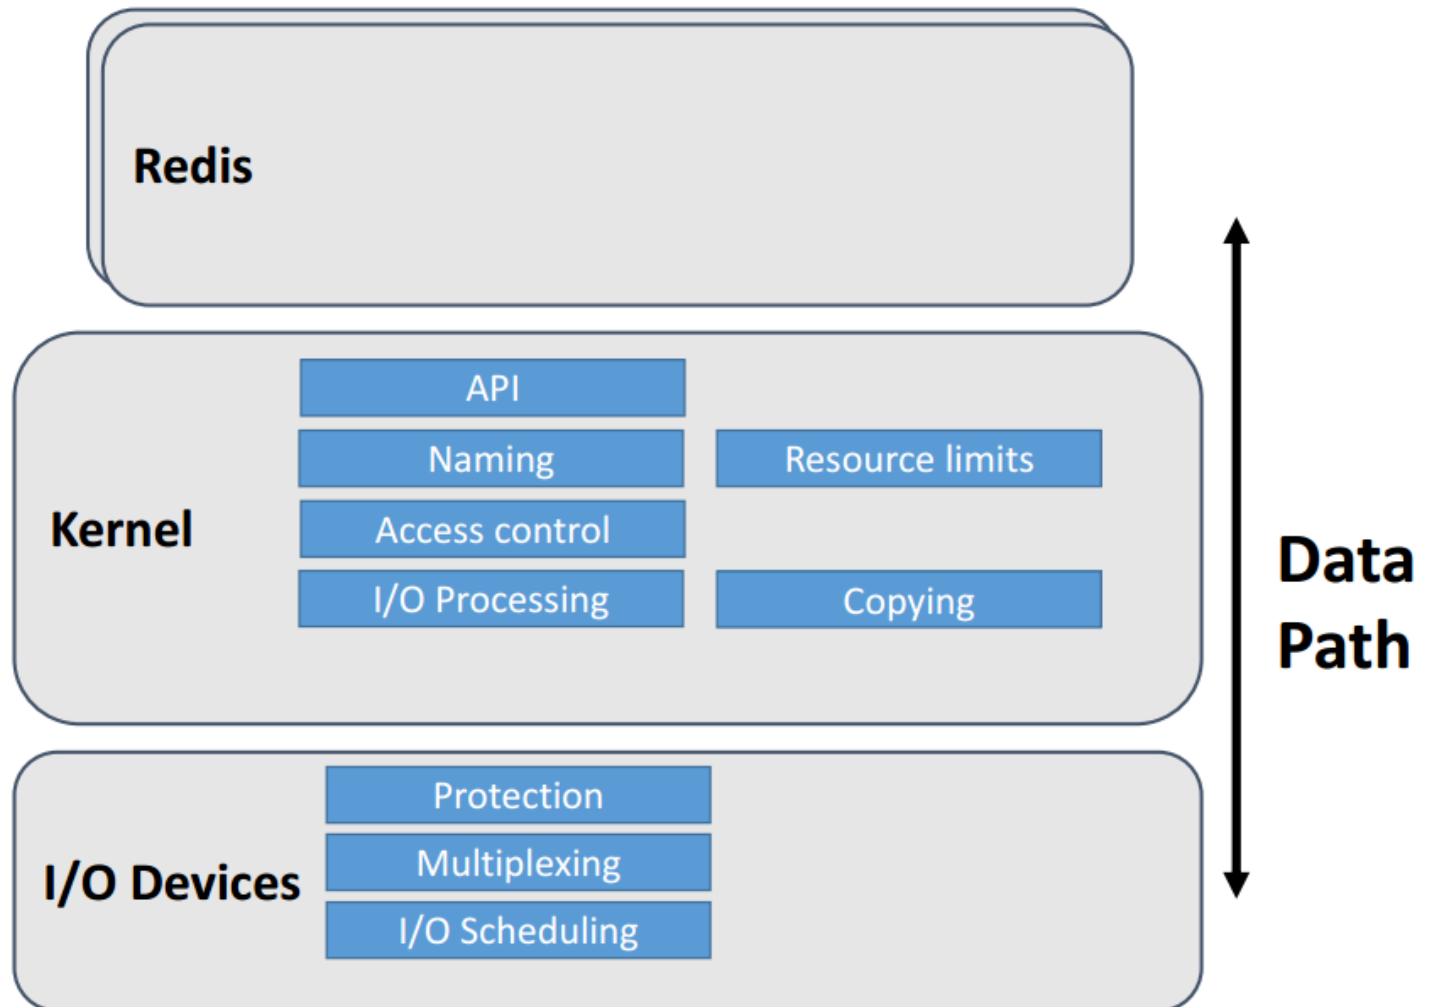
- Standard on NIC, emerging on RAID
- Multiplexing
 - SR-IOV: Virtual PCI devices w/ own registers, queues, INTs
- Protection
 - IOMMU:
Devices use app virtual memory
 - Packet filters, logical disks:
Only allow eligible I/O
- I/O Scheduling
 - NIC rate limiter, packet schedulers



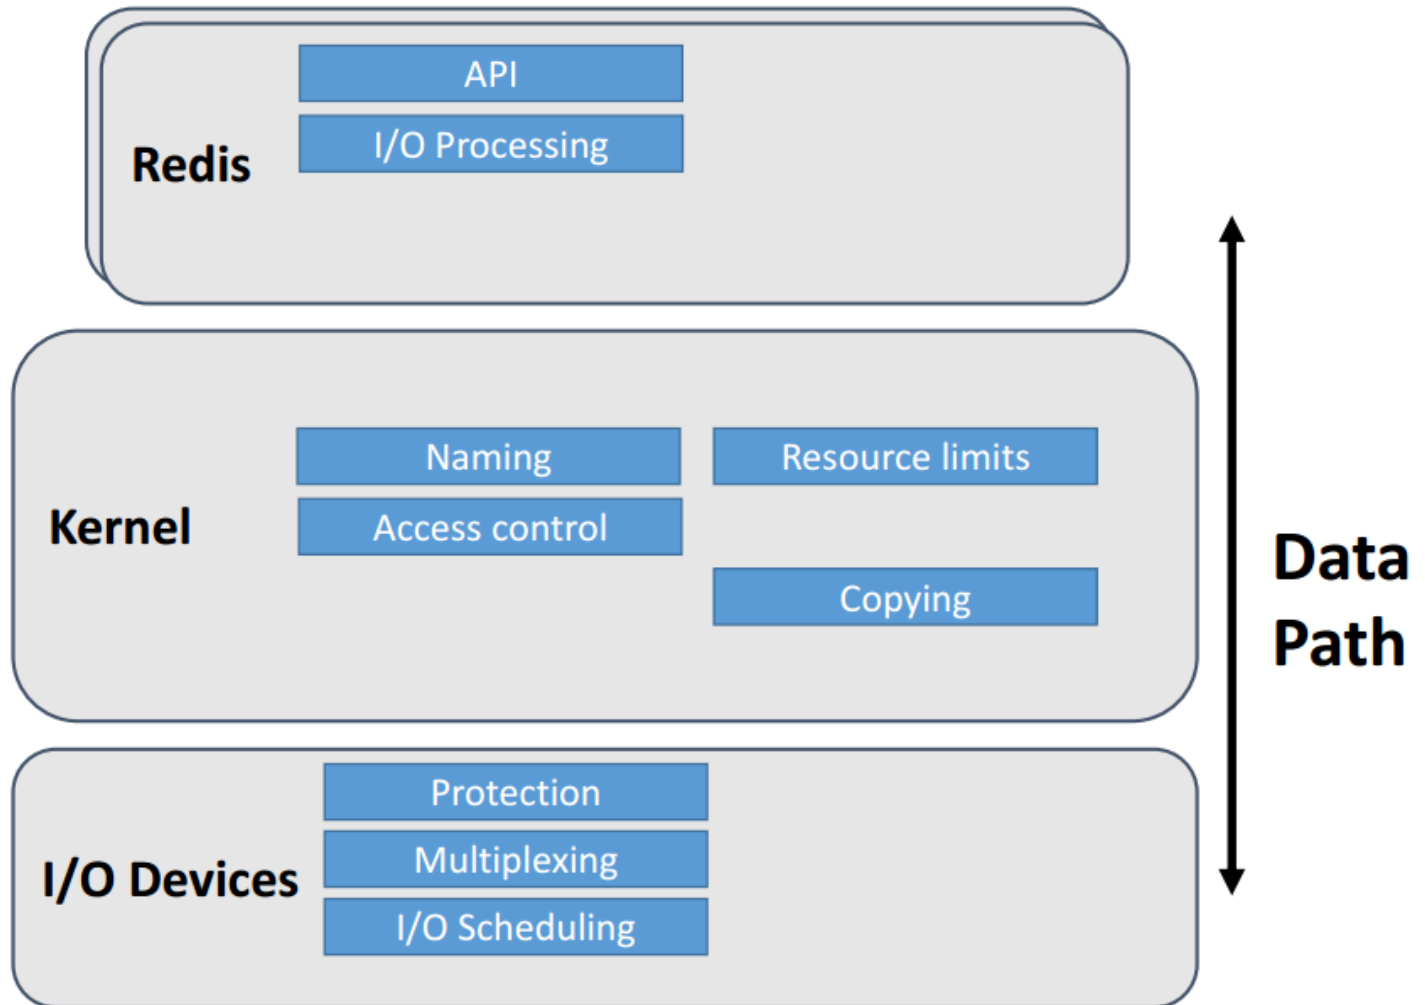
How to skip the kernel?



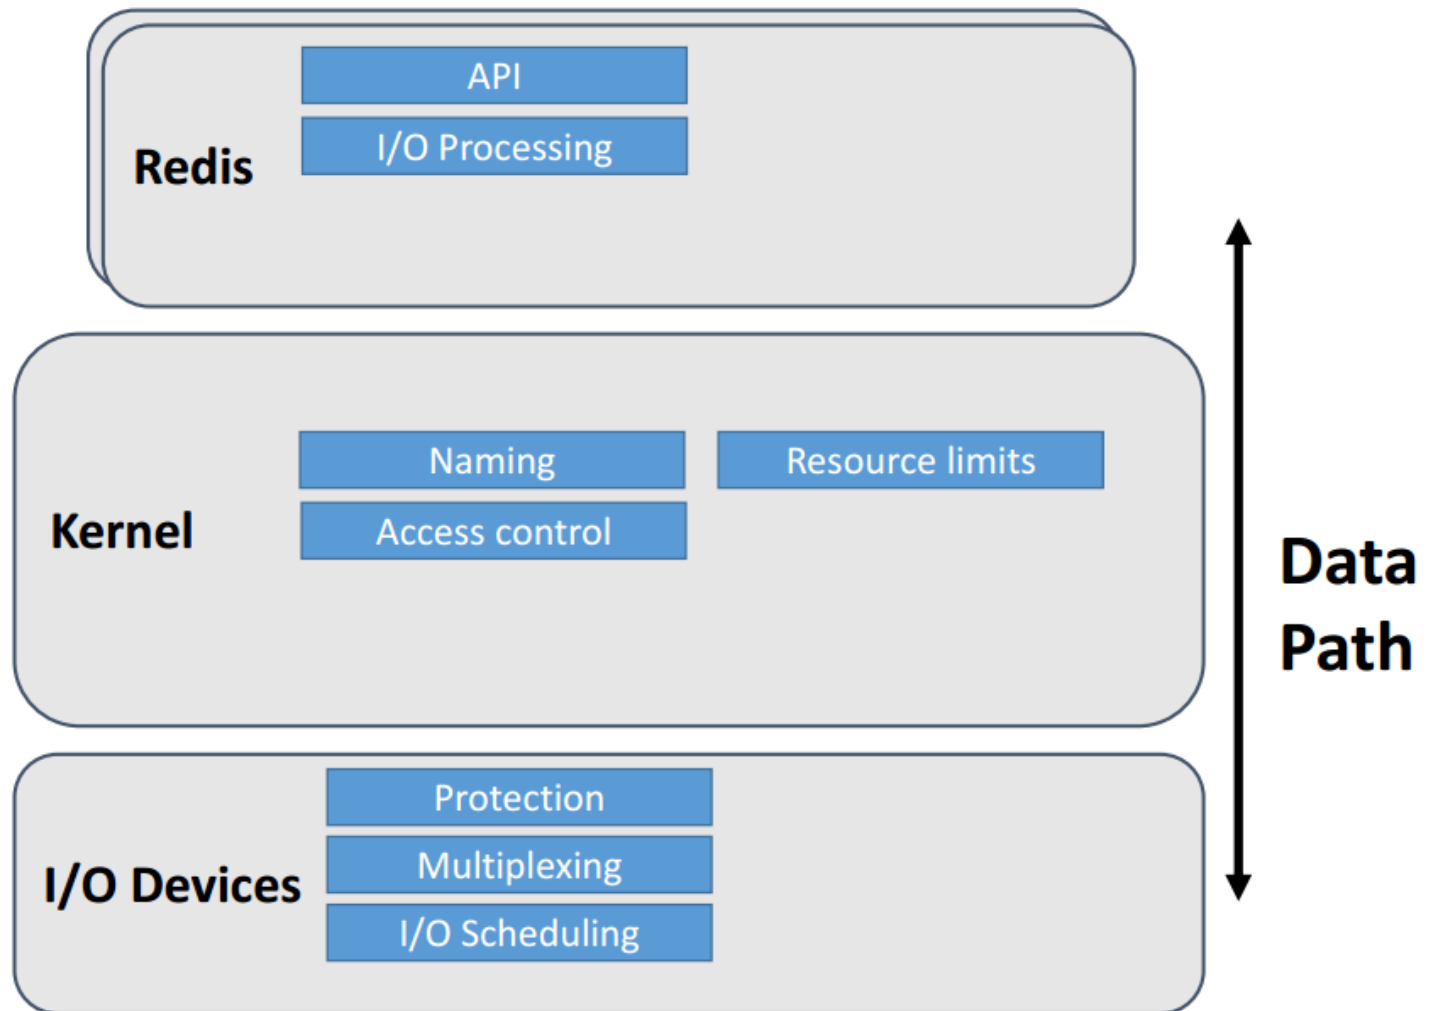
How to skip the kernel?



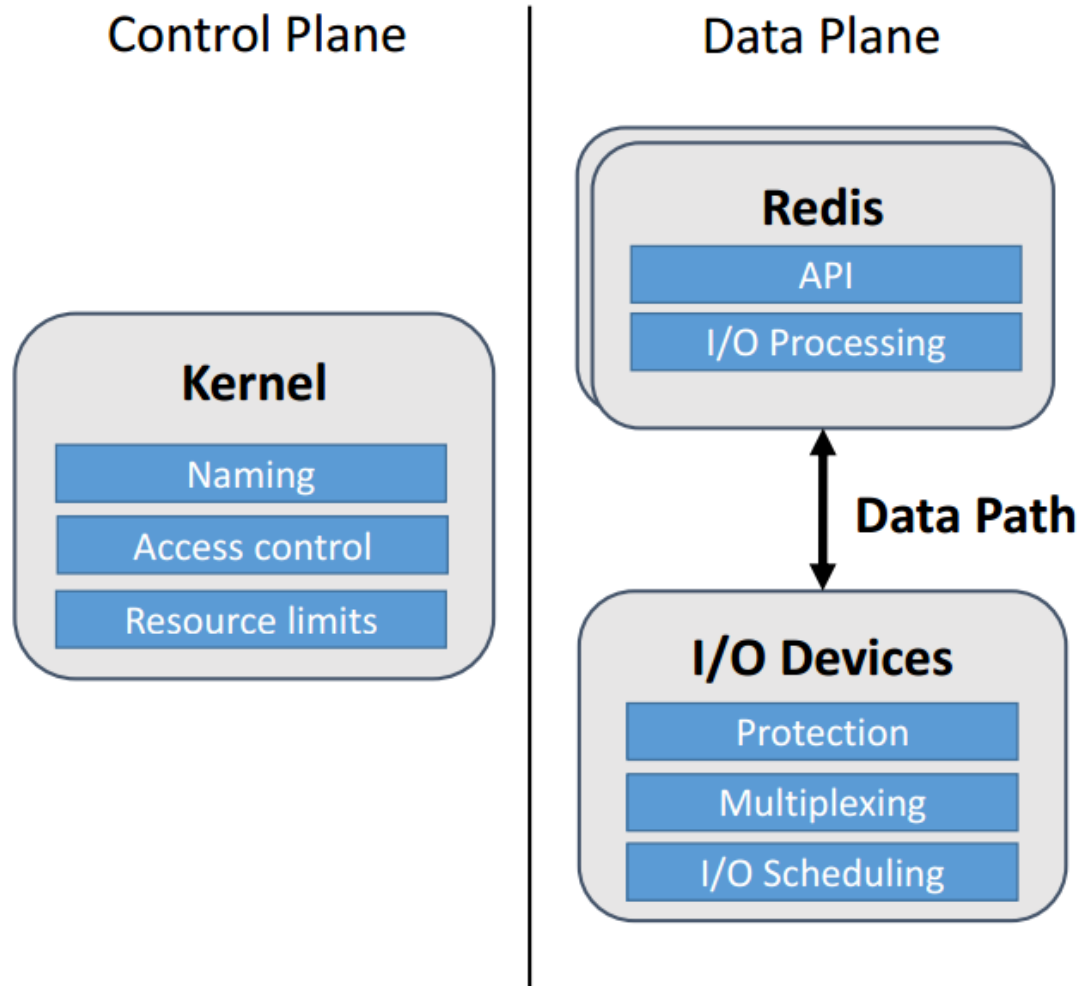
How to skip the kernel?



How to skip the kernel?



Arrakis I/O Architecture



Arrakis I/O Architecture





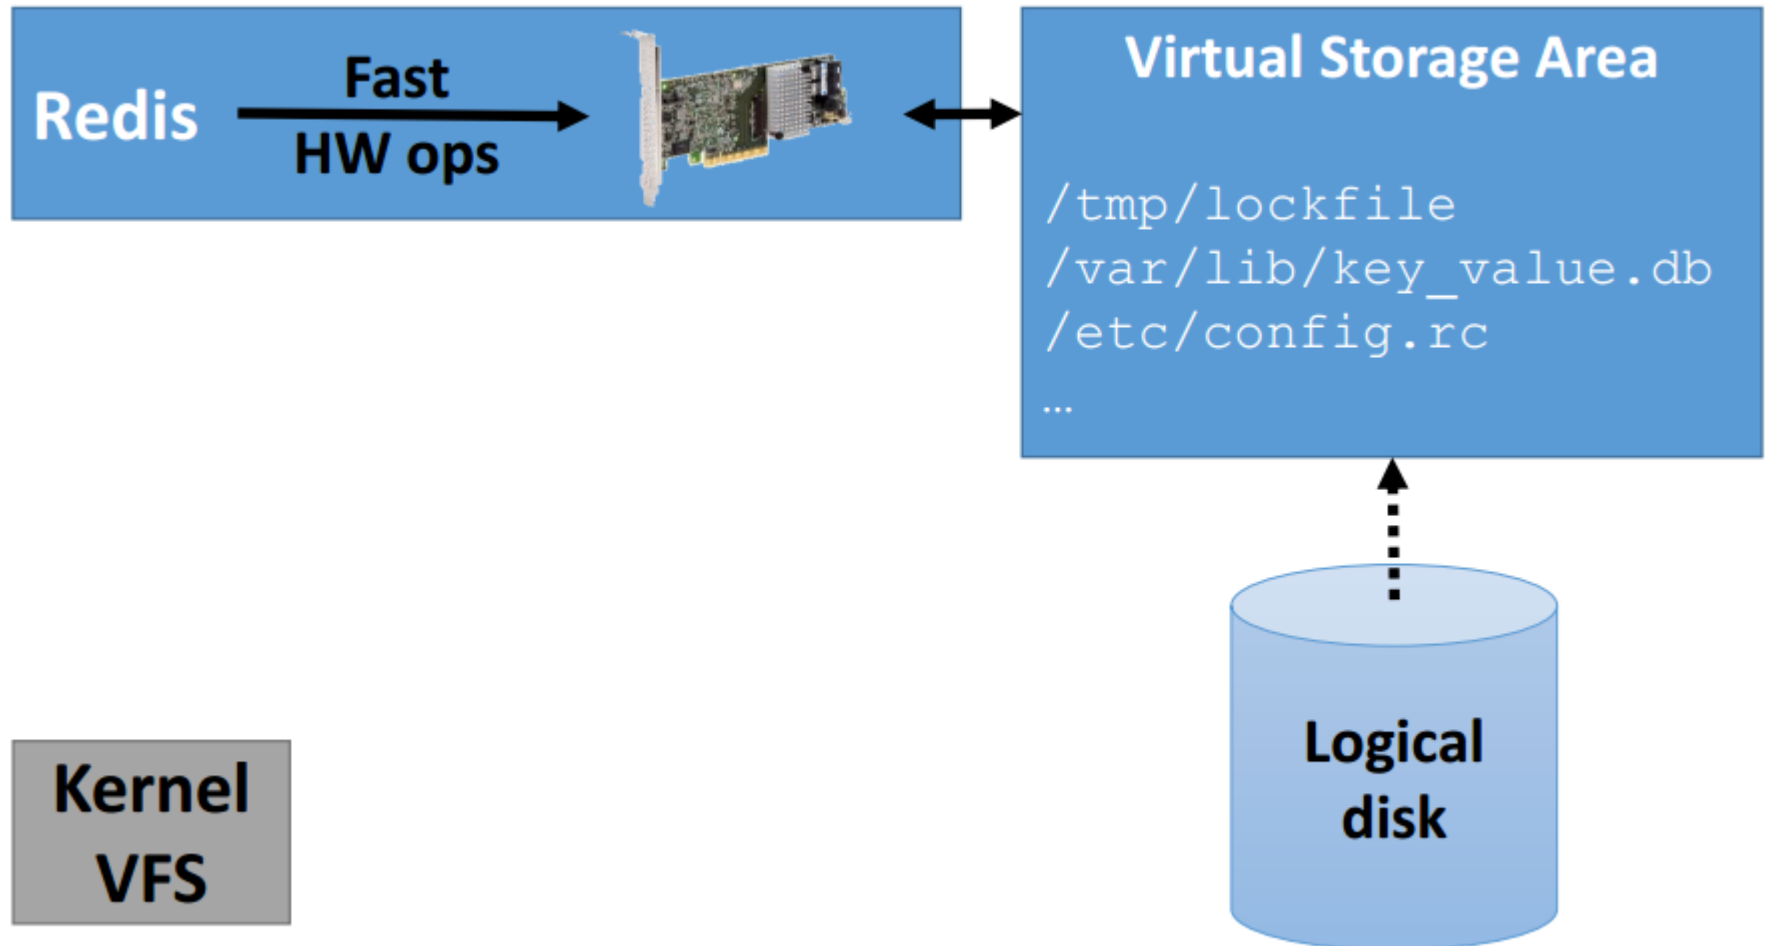
Fig. 4. Arrakis architecture. The storage controller maps VSAs to physical storage.

Arrakis Control Plane

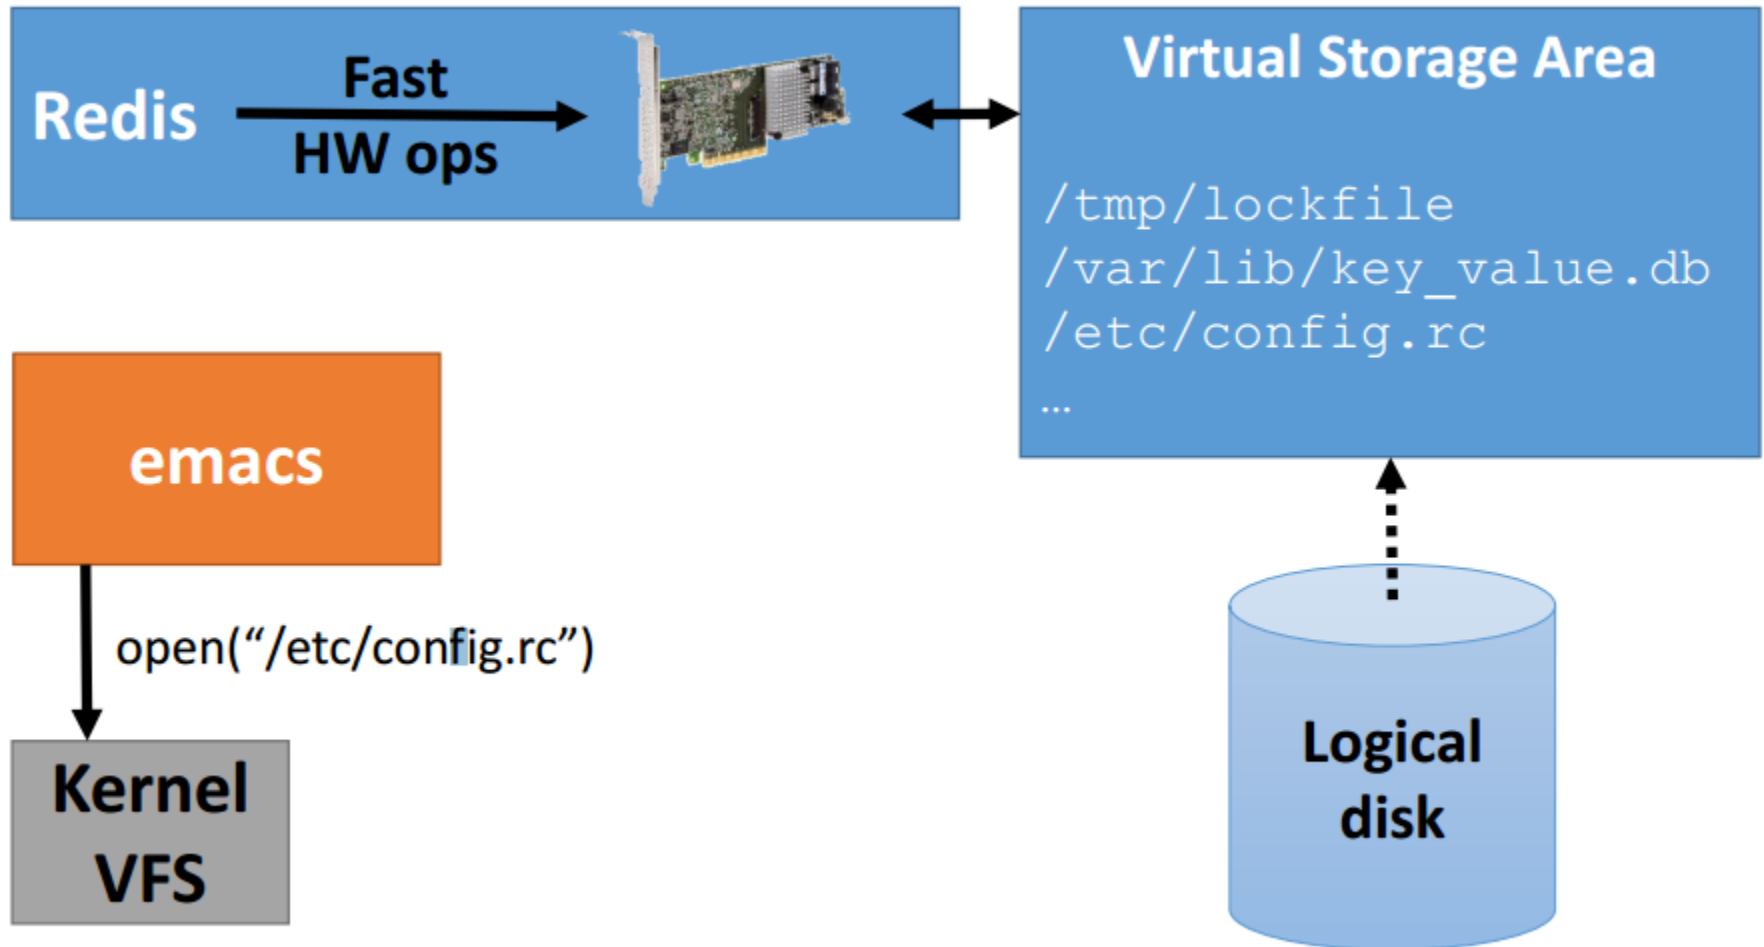
- Access control
 - Do once when configuring data plane
 - Enforced via NIC filters, logical disks
- Resource limits
 - Program hardware I/O schedulers
- Global naming
 - Virtual file system still in kernel
 - Storage implementation in applications



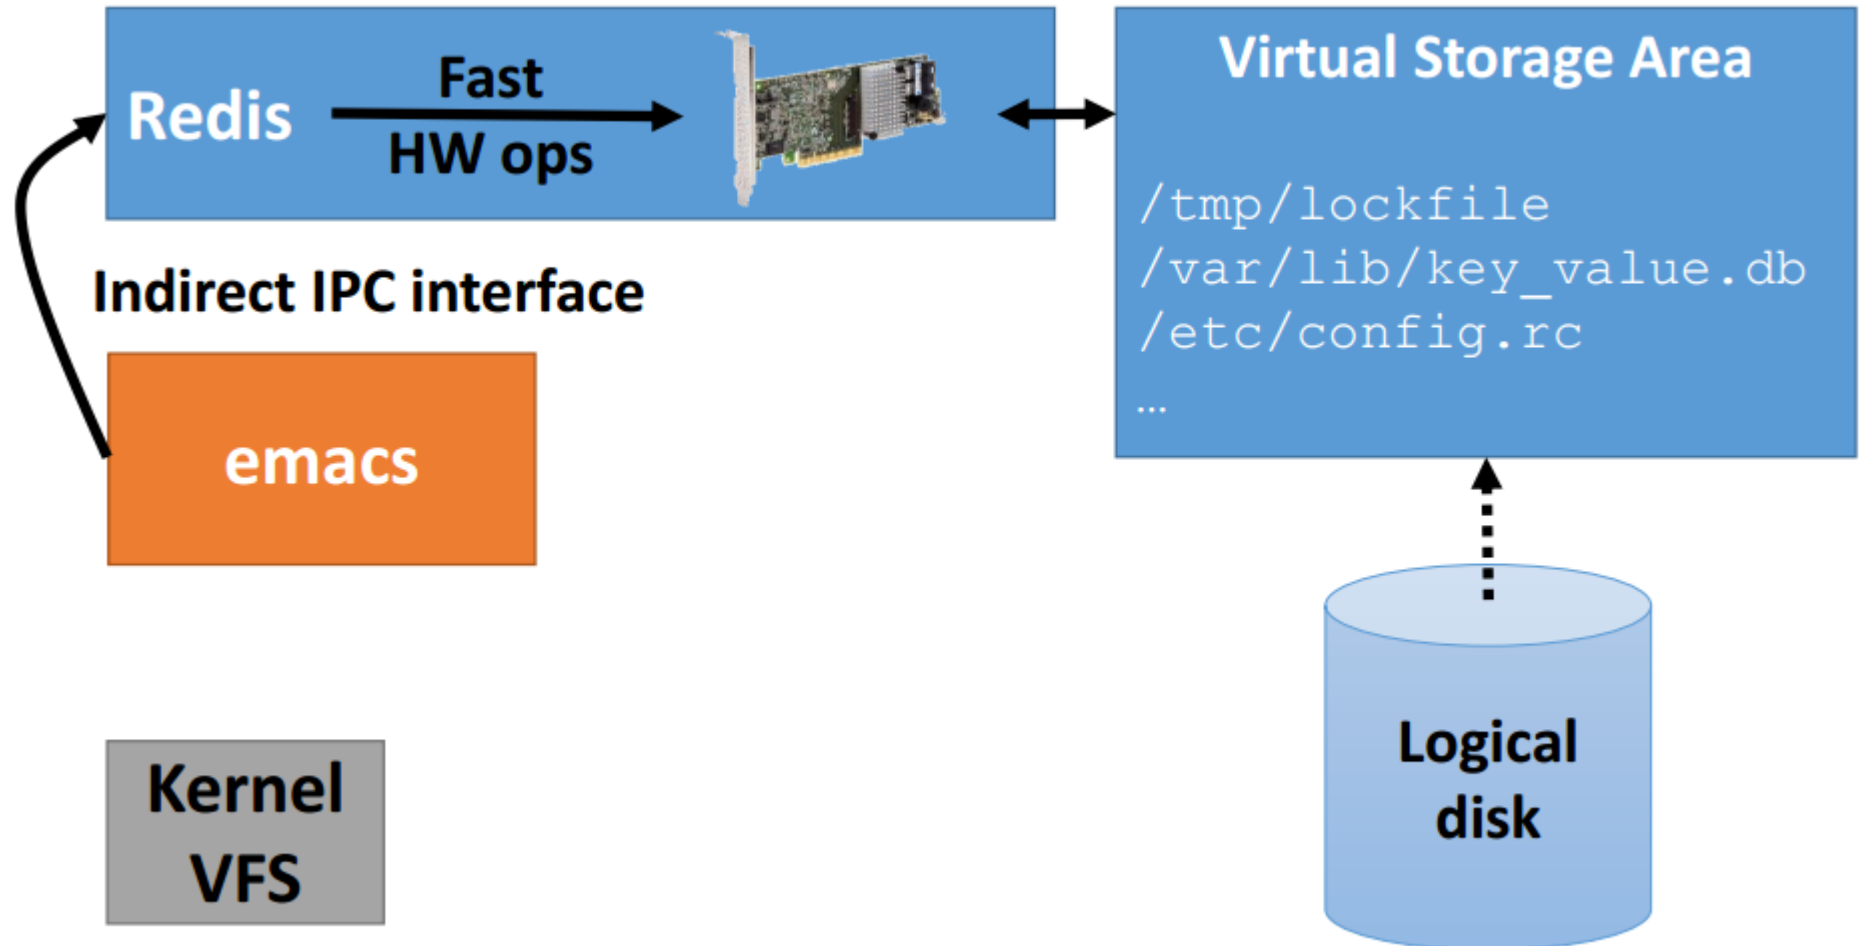
Global Naming



Global Naming




Global Naming



Storage Data Plane: Persistent Data Structures

- Examples: **log, queue**
- Operations immediately persistent on disk




Benefits:


- In-memory = on-disk layout
 - Eliminates marshaling
- Metadata in data structure
 - Early allocation
 - Spatial locality
- Data structure specific caching/prefetching
- Modified Redis to use **persistent log: 109 LOC** changed

Evaluation: Redis Latency

- Reduced (in-memory) GET latency by **65%**



- Reduced (persistent) SET latency by **81%**




Evaluation:


Redis Throughput

- Improved GET throughput by **1.75x**
 - Linux: **143k** transactions/s
 - Arrakis: **250k** transactions/s
- Improved SET throughput by **9x**
 - Linux: **7k** transactions/s
 - Arrakis: **63k** transactions/s

Evaluation: memcached Scalability



Evaluation: Single-core Performance



Summary

- OS is becoming an I/O bottleneck
 - Globally shared I/O stacks are slow on data path
- **Arrakis:** Split OS into control/data plane
 - Direct application I/O on data path
 - Specialized I/O libraries
- Application-level I/O stacks deliver great performance
 - **Redis:** up to **9x** throughput, **81%** speedup
 - Memcached **scales linearly** to **3x** throughput